



Support Vector Machines


Maximizing the Margin



Anoop M. Namboodiri
Centre for Visual Information Technology
IIIT, Hyderabad, INDIA



Perceptron Learning

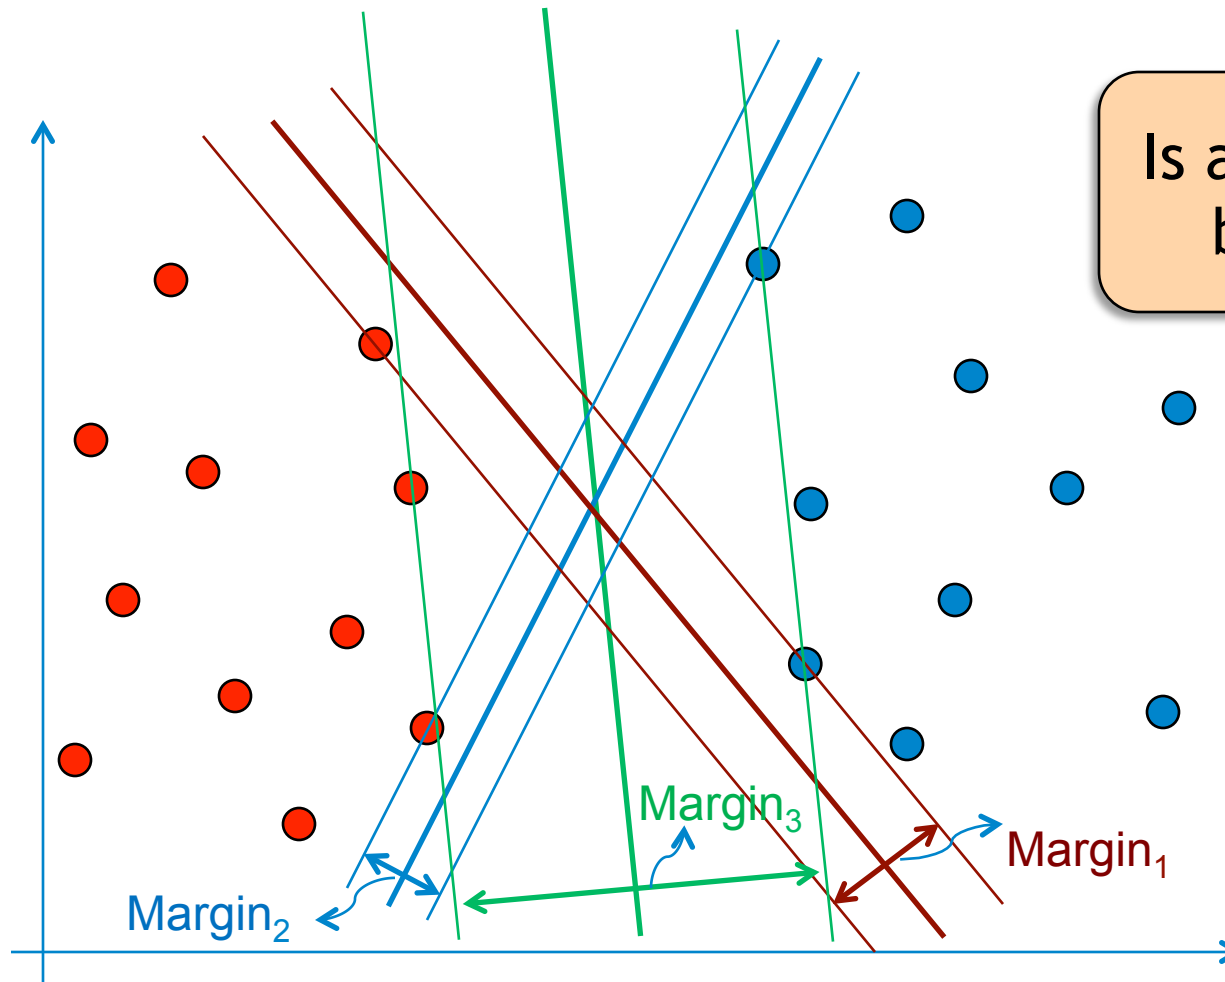


Are all solutions equally good?

- Multiple solutions exist for linearly separable data
- Perceptron learning (any GD) results in a feasible solution




Margin: The No-mans Band



Is a Larger Margin better? Why?

- Margin: Width of a band around decision boundary without any training samples
- Margin varies with the position and orientation of the separating hyperplane

Margin: Bubbles around Samples




A few samples control the Dec. boundary

- Margin: Radius of a region around each training sample, through which the decision boundary cannot pass
- As margin increases, the feasible region reduces



Margin: Band vs. Bubbles




Samples that support the Decision boundary are called *Support Vectors*

- Margin: Both interpretations yield the same decision boundary



Can we Quantify Generalization?

- Work by Vapnik and Chervonenkis in the 1970s
 1. V.N. Vapnik, A.Ya. Červonenkis, “On Uniform Frequencies of Events to their Probabilities”, *Primenen*, 1971, 16(2), pp. 264-279.
 2. V.N. Vapnik, “Estimation of Dependences Based on the Empirical Data” [Russian]. Nauka, Moscow, 1979.
 3. V.N. Vapnik, “The Nature of Statistical Learning Theory”. Springer-Verlag, New York, 1995.



- Bound on Expected loss [3]:

$$R(\alpha) \leq R_{train}(\alpha) + \sqrt{\frac{f(h)}{N}}$$

- h is the VC dimension, and $f(h)$ is given by:

$$f(h) = h + h \log(2N) - h \log(h) - c$$




Why Maximize the Margin?

- To reduce test error, keep training error low (say 0), and minimize the VC-dimension, h .

Relative Margin : ρ/D

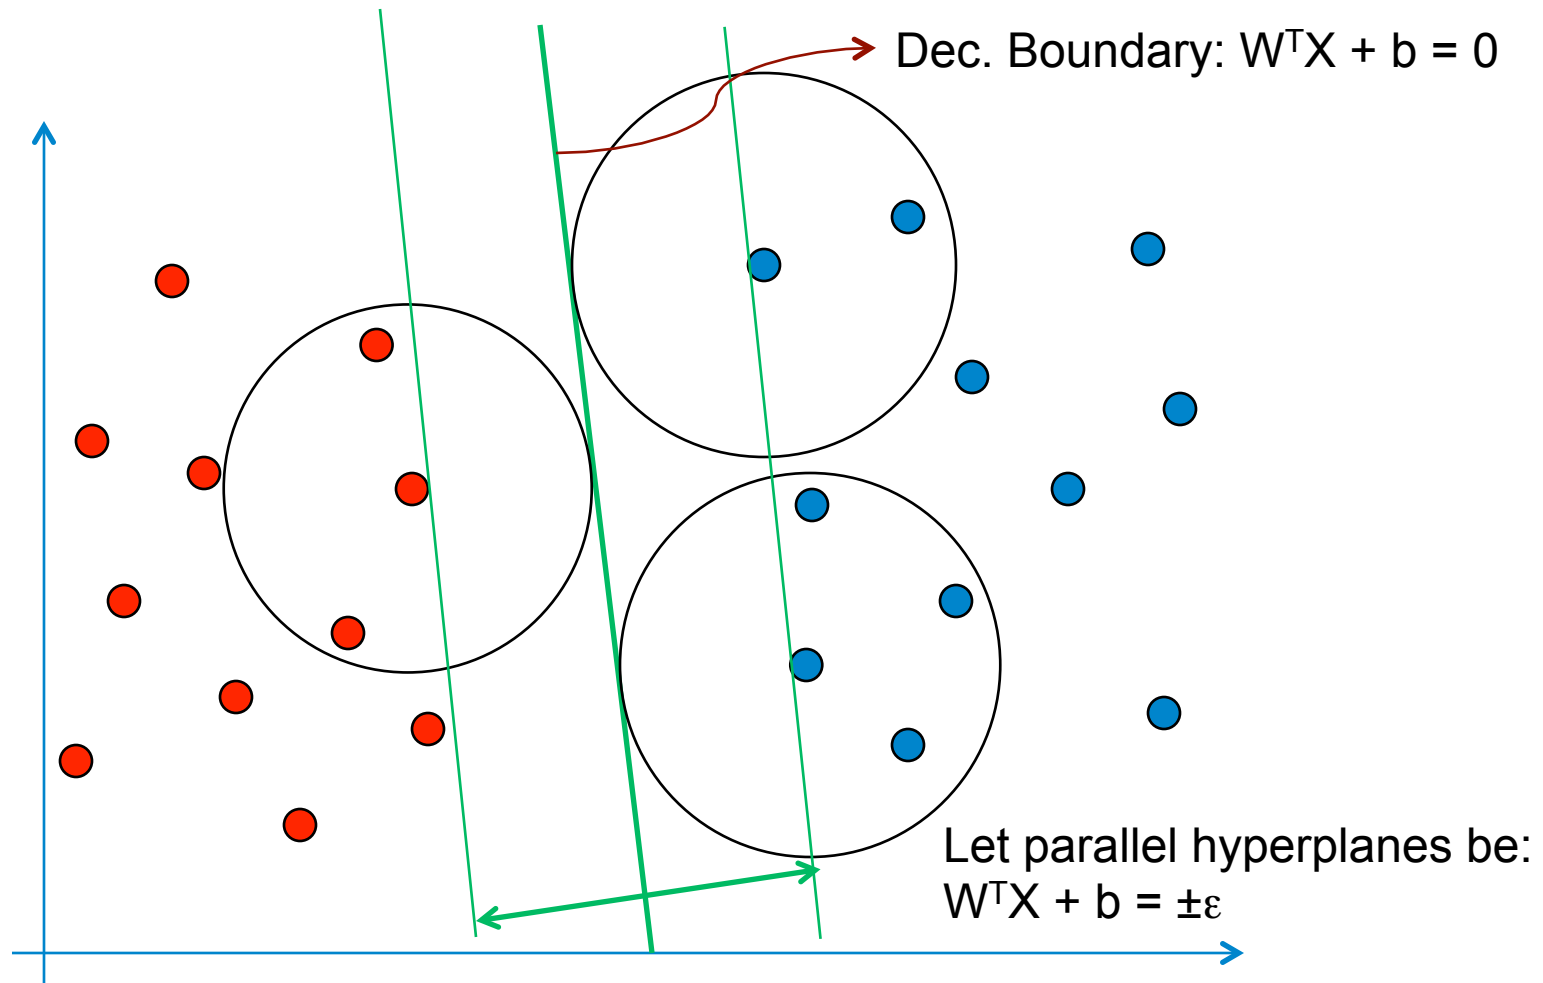
$$\text{VC-D, } h \leq \min \left\{ d, \left\lceil \frac{D^2}{\rho^2} \right\rceil \right\} + 1$$

- Maximizing margin improves generalization.
- h can be made independent of the dimensionality: d .





Formalizing the Margin




- Note: The value of $W^T X_i + b$ is dependent on the scale of X and W



Formulation

- Let $g(x) = w^T x + b$.
- We want to maximize k such that:
 - $w^T x_i + b \geq k$ for $d_i = 1$
 - $w^T x_i + b \leq -k$ for $d_i = -1$
- Value of $g(x)$ depends on $\|w\|$:
 1. Keep $\|w\| = 1$, and maximize $g(x)$, or
 2. Let $g(x) \geq 1$, and minimize $\|w\|$.
- We use approach (2) and formulate the problem as:
 - Minimize: $\frac{1}{2} w^T w$
 - Subject to: $d_i(w^T x_i + b) \geq 1$, for $i = 1..N$





The Optimization Problem


$$\text{Minimize: } \Phi(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w}$$

$$\text{Subject to: } d_i(\mathbf{w}^T \mathbf{x}_i + b) - 1 \geq 0 \quad \forall i$$

- Quadratic objective function with linear inequalities as constraints: QP Solver.
- Integrating the constraints into the Lagrangian form, we get:

$$\text{Minimize: } J(\mathbf{w}, b, \boldsymbol{\alpha}) = \frac{1}{2} \mathbf{w}^T \mathbf{w} - \sum_{i=1}^N \alpha_i d_i(\mathbf{w}^T \mathbf{x}_i + b) + \sum_{i=1}^N \alpha_i$$

$$\text{Subject to: } \alpha_i \geq 0 \quad \forall i$$



- Minimize J with respect to \mathbf{w} and b , and maximize with respect to $\boldsymbol{\alpha}$.

Converting to the Dual Form



$$\text{Objective: } J(\mathbf{w}, b, \boldsymbol{\alpha}) = \frac{1}{2} \mathbf{w}^T \mathbf{w} - \sum_{i=1}^N \alpha_i d_i (\mathbf{w}^T \mathbf{x}_i + b) + \sum_{i=1}^N \alpha_i$$

At the optimum,

$$1: \frac{\partial J}{\partial \mathbf{w}} = 0$$

$$\text{and } 2: \frac{\partial J}{\partial b} = 0$$

KKT Conditions

$$1: \mathbf{w}_o = \sum_{i=1}^N \alpha_i d_i \mathbf{x}_i$$

$$2: \sum_{i=1}^N \alpha_i d_i = 0$$

$$3: \alpha_i [d_i (\mathbf{w}_o^T \mathbf{x}_i + b_o) - 1] = 0$$

$$\text{Obj: } J(\mathbf{w}, b, \boldsymbol{\alpha}) = \sum_{i=1}^N \alpha_i + \frac{1}{2} \mathbf{w}^T \mathbf{w} - \mathbf{w}^T \sum_{i=1}^N \alpha_i d_i \mathbf{x}_i - b \sum_{i=1}^N \alpha_i d_i$$

$$\text{Using 1, 2: } Q(\boldsymbol{\alpha}) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j d_i d_j \mathbf{x}_i^T \mathbf{x}_j$$



Solving the Dual Form

$$Q(\alpha) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j d_i d_j \mathbf{x}_i^T \mathbf{x}_j$$

$$\text{Subject to } \alpha_i \geq 0 \quad \forall_i \quad \text{and} \quad \sum_{i=1}^N \alpha_i d_i = 0$$

QP Solver

α_i

$$\mathbf{w}_o = \sum_{i=1}^N \alpha_i d_i \mathbf{x}_i$$

KKT Conditions

$$\alpha_i [d_i (\mathbf{w}_o^T \mathbf{x}_i + b_o) - 1] = 0$$



$$b_o = 1 - \mathbf{w}_o^T \mathbf{x}_{s+}$$

- The only unknowns (variables) are α_i s.
- The constraints are also on α_i s only.
- Data vectors appear only as dot products
- Objective is convex, subject to linear constraints
- Can be solved using standard convex quadratic program solvers

Non-Separable Data

1: Noisy Data/Bad Features

2: Non-linear Boundary





Noise in Data

Minimize : $\Phi(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w}$

Subject to : $d_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 \quad \forall i$

Introduce slack variables $\xi_i \geq 0$


Minimize : $\Phi(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w}$

Subject to : $d_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \xi_i \quad \forall i$

Also minimize training error $\sum_{i=1}^N I(\xi_i \geq 1)$ or $\sum_{i=1}^N \xi_i$

Minimize : $\Phi(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum_{i=1}^N \xi_i$

Subject to : $d_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \xi_i ; \quad \xi_i \geq 0, \quad \forall i$





Dual form with Slack


- Forming the Lagrangian and converting to dual, we get:

$$Q(\alpha) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j d_i d_j \mathbf{x}_i^T \mathbf{x}_j$$

$$\text{Subject to } 0 \leq \alpha_i \leq C \quad \forall_i \quad \text{and} \quad \sum_{i=1}^N \alpha_i d_i = 0$$

- Note that neither the slack variables, nor their Lagrange multipliers appear in the dual.
- The only change is the additional constraint on α_i
- The parameter C controls the relative weight between training error and the VC dimension.

Non-Linear Boundaries



Φ is a non-linear mapping into a possibly high-dimensional space



Solving the SVM with a Mapping

- Data vectors occur only as dot products in SVM-learning and testing

Testing Phase

- Label = $\text{sign}(\mathbf{w}_o \cdot \Phi(\mathbf{x}_{\text{test}}) + b_o)$

$$\mathbf{w}_o = \sum_{i=1}^N \alpha_i d_i \Phi(\mathbf{x}_i)$$

$$\therefore L \quad Q(\boldsymbol{\alpha}) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j d_i d_j K(\mathbf{x}_i, \mathbf{x}_j)$$

$$\text{Label} = \text{sign} \left(\sum_{i=1}^N (\alpha_i d_i K(\mathbf{x}_i, \mathbf{x}_{\text{test}})) + b_o \right)$$



A Simple Quadratic Kernel

$$\text{Let } \Phi(\mathbf{X}) = \Phi\left(\begin{bmatrix} x_1 \\ x_2 \end{bmatrix}\right) = \begin{bmatrix} x_1 x_1 \\ x_1 x_2 \\ x_2 x_1 \\ x_2 x_2 \end{bmatrix}$$

We can compute $K(\mathbf{X}, \mathbf{Y}) = (\mathbf{X} \cdot \mathbf{Y})^2$ instead of mapping with Φ explicitly and then computing dot product.

$$\text{Let } K(\mathbf{X}, \mathbf{Y}) = \Phi(\mathbf{X}) \cdot \Phi(\mathbf{Y}) = \begin{bmatrix} x_1^2 \\ x_1 x_2 \\ x_2 x_1 \\ x_2^2 \end{bmatrix} \cdot \begin{bmatrix} y_1^2 \\ y_1 y_2 \\ y_2 y_1 \\ y_2^2 \end{bmatrix}$$

$$= x_1^2 y_1^2 + 2x_1 x_2 y_1 y_2 + x_2^2 y_2^2 = (x_1 y_1 + x_2 y_2)^2 = (\mathbf{X} \cdot \mathbf{Y})^2$$



Similarly for a Cubic Kernel

- Original Space: 2-dimensional

$$\text{Let } K(\mathbf{X}, \mathbf{Y}) = (\mathbf{X} \cdot \mathbf{Y})^3 = (x_1 y_1 + x_2 y_2)^3$$

$$\Phi(\mathbf{X}) = \Phi\left(\begin{bmatrix} x_1 \\ x_2 \end{bmatrix}\right) = \begin{bmatrix} x_1^3 \\ x_1^2 x_2 \\ x_1 x_2^2 \\ x_2^3 \end{bmatrix}$$

- Equivalent to working in an 4-dimensional space
- 4:× and 1:+, instead of 16:× and 3:+



Similarly for a Cubic Kernel

- Original Space: 3-dimensional

$$\text{Let } K(\mathbf{X}, \mathbf{Y}) = (\mathbf{X} \cdot \mathbf{Y})^3 = (x_1 y_1 + x_2 y_2 + x_3 y_3)^3$$

$$\Phi(\mathbf{X}) = \Phi\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}\right) = \begin{bmatrix} x_1^3 \\ x_2^3 \\ x_3^3 \\ x_1^2 x_2 \\ x_1 x_2^2 \\ x_1^2 x_3 \\ x_1 x_3^2 \\ x_2^2 x_3 \\ x_2 x_3^2 \\ x_1 x_2 x_3 \end{bmatrix}$$

- Equivalent to working in an 10-dimensional space
- 5:× and 2:+, instead of 38:× and 9:+



A Generic Polynomial Kernel

- Adding two Kernels gives you a new Kernel:

$$K(\mathbf{X}, \mathbf{Y}) = K_1(\mathbf{X}, \mathbf{Y}) + K_2(\mathbf{X}, \mathbf{Y}) \quad : \quad \Phi(\mathbf{X}) = \begin{bmatrix} \Phi_1(\mathbf{X}) \\ \Phi_2(\mathbf{X}) \end{bmatrix}$$

$$K_p(\mathbf{X}, \mathbf{Y}) = (1 + \mathbf{X} \cdot \mathbf{Y})^p = 1 + \mathbf{X}^1 \cdot \mathbf{Y} + (\mathbf{X} \cdot \mathbf{Y})^2 + \dots + (\mathbf{X}^2 \cdot \mathbf{Y})^p$$

- Just adding a 1 and raising to the power of p maps the input vector into a space containing all original dimensions, all 2-products, 3-products, ..., p-products.

Mercer's Theorem

- Using Kernels, we avoid explicit mapping with Φ
- In fact, we do not even have to know what Φ is as long as we are sure there exists a valid Φ .
- Mercer's Theorem:

Any given kernel can be expanded as a series :

$$K(\mathbf{X}, \mathbf{Y}) = \sum_{i=1}^{\infty} \lambda_i \Phi_i(\mathbf{X}) \cdot \Phi_i(\mathbf{Y}), \quad \lambda_i > 0; \text{ iff}$$

$K(\mathbf{X}, \mathbf{Y})$ satisfies the Mercer's conditions
(symmetric, continuous, positive semi - definite)

Popular Kernels

- Polynomial:

$$K_p(\mathbf{X}, \mathbf{Y}) = (1 + \mathbf{X} \cdot \mathbf{Y})^p$$

- Radial Basis Function (RBF) or Gaussian:

$$K_r(\mathbf{X}, \mathbf{Y}) = e^{-\frac{1}{2\sigma^2} \|\mathbf{X} - \mathbf{Y}\|_2^2}$$

- Hyperbolic Tangent:

$$K_s(\mathbf{X}, \mathbf{Y}) = \tanh(\beta_0 \mathbf{X} \cdot \mathbf{Y} + \beta_1)$$