


CSE578: Computer Vision

Spring'16

Multiple-View Structure Recovery




Anoop M. Namboodiri

Center for Visual Information Technology

IIIT Hyderabad, INDIA

Multiple Views of Points/Objects

- Given projections of a set of 3D points in two or more cameras, get their 3D coordinates.
- Each 3D point is identified in every camera view.
- What else is known? Camera matrices K_i , R_i , t_i ? Only the intrinsic parameters K_i ?



Variations of the Problem

- **(Binocular) Stereo:** Two cameras with known intrinsic and extrinsic parameters.
- **Multiview Stereo:** Multiple known cameras
- **Structure-from-Motion:** Given m cameras and n points and projections \mathbf{x}_{ij} of point j in camera i , recover 3D points \mathbf{X}_j and camera matrices \mathbf{C}_i
 - **Affine SfM:** For affine cameras
 - **Projective SfM:** For general projective cameras
- **Bundle Adjustment:** Directly recover $\mathbf{C}_i, \mathbf{X}_j$ from \mathbf{x}_{ij}


Binocular Stereo and Feature Correspondence

Geometry of Matching

- B : baseline, f : focal length, z : depth, x : image coords
- From similar triangles:

$$\frac{x_1 - x_2}{f} = \frac{B}{z}$$

- **Stereo Disparity or Parallax:** the “shift” between the left and right images. $\Delta = \frac{Bf}{z}$.




- Farther the point, smaller the disparity and vice versa
- A large baseline can give more reliable estimates of depth. But, matching may become harder
- Basic step: **Identify common points in camera views**

Identifying Common Points

- Find a world point in 2 or more views
- Appearance is the only clue to identify them
- Individual pixel colours are similar very often. Match is too noisy
- Match a (small) neighbourhood of colours from one image to a similar neighbourhood in others
- Will work if local surface is fronto-parallel and images have similar magnification
- Foreshortening can happen when viewing an oblique surface
- Many ambiguities. We need a lot of help!

Some Examples



Matching Patches

- Compare $m \times m$ patches from two views.
Form vectors \mathbf{v} and \mathbf{v}' of length m^2 from them
- Matching scores between patches:
 - Sum of Absolute Difference (SAD): $\|\mathbf{v} - \mathbf{v}'\|_1$
 - Sum of squared difference (SSD): $\|\mathbf{v} - \mathbf{v}'\|_2$
 - Correlation: $\frac{\mathbf{v}'^T \mathbf{v}}{\sqrt{\mathbf{v}^T \mathbf{v}} \sqrt{\mathbf{v}'^T \mathbf{v}'}}$
 - Normalized correlation: $\frac{\bar{\mathbf{v}}^T \bar{\mathbf{v}'}}{\sqrt{\bar{\mathbf{v}}^T \bar{\mathbf{v}}} \sqrt{\bar{\mathbf{v}}'^T \bar{\mathbf{v}}'}} \cdot \text{Range: } [-1, 1]$
 $\bar{\mathbf{v}}, \bar{\mathbf{v}}'$ are vectors with respective patch-mean colour subtracted.
 - Invariant to affine changes in intensity/colour.


Constraints on Matching

- **Epipolar:** Match lies on the epipolar line of the pixel
- **Colour Constancy:** The appearance/colour does not change from one view to another
- **Uniqueness:** A point on left image can match with only one point on the right and vice versa
- **Ordering or Monotonicity:** If point A is to the left of B in view 1, it will be to the left of B in view 2 also. (Violated if great difference in depth exists)
- **Continuity:** Disparity values vary smoothly (violated at occlusion boundaries)

Sparse correspondence: only for good feature points

Dense correspondence: a match for every pixel


Epipolar Geometry





Reduced Search and Rectification

- The search is limited to a line if fundamental matrix known (i.e., **weakly calibrated**)
- Simplest if left and right cameras have same image plane and pure X-translation between them.
- Fundamental matrix has a simple form. Epipolar constraint reduces to $y' = y$.
- Matches constrained to lie in the same scan line
- **Rectification:** A rotation of the camera (to make image planes parallel) and a change in **K** matrix (focal length, image center).
- Can be represented using a homography **H** to align one image plane to the other
Or, homographies **H₁**, **H₂** to align them to a third plane



Rectification



Rectification: Example



Ordering Constraint



Order of matches from left and right is ordinarily preserved
Violation may mean something else.

Various Situations




Image courtesy Marc Pollefeys

Shows the *epipolar line image* or *disparity space image* with different scenarios


Scan-Line to Scan-Line Matching

- Disparity Line Image pits pixels of one row of the left image against the pixels of the matching epipolar line in the other.
- Several matching scenarios:
 - If left pixel $(i - 1)$ matches with right pixel $(j - 1)$, next pixel i can match pixel j , if the match is good
 - Otherwise, it may continue the match with $(j - 1)$ with an occlusion cost (due to left occlusion)
 - Or, $(i - 1)$ can match with j with another occlusion cost (due to right occlusion)
- Sub-pixel definitions may be needed when zoom is different

Dynamic Programming Solution

- Cost of matching: $C(i-1, j-1) + c(i, j)$ if pixels match, $C(i-1, j) + C_o$ if left occlusion, and $C(i, j-1) + C_o$ if right occlusion, where C_o is a high occlusion cost
- Select the minimum from those three and declare match or occlusion accordingly
- Can be setup nicely as a dynamic programming solution working in the i, j space, starting with leftmost pixel match
- Cost of matching: $O(N^2)$ where N is the number of pixels in each scanline.

Dynamic Programming Path



- Initialize first row and col to $i * C_o$
- Do for $i \in [0, N - 1]$ and $j \in [0, N - 1]$:
Set $C(i, j)$ to min of $C(i - 1, j - 1) + c(i, j)$, $C(i - 1, j) + C_o$, $C(i, j - 1) + C_o$
- Mark each as **M/L/R**
- Reconstruct from (N, N) , by following the **M** pixels and their connections.


Globally Optimal Solution

- Provides a *globally optimal* match as opposed to the local matching done by search
- Provides dense correspondence: a match for every pixel
- Works well enough. And is a prototype for many global stereo matching approaches that followed
- Difficulty: assigning a cost for occlusions.
- Difficulty: maintaining consistency across scan lines

We will see another global solution using graphcuts later

Structured Lighting

- Finding correspondences is hard by itself
- Can we help it by projecting patterns onto the world?
Structured Lights!
- Lightstrip range finders, etc.
- Combination of sinusoids sometimes to get dense matches
- *Active vision*, as it changes the appearance
- The light projected need not be in the visible spectrum



Xbox Kinect

IR-based range sensor for Xbox

- Aligned depth and RGB images at 640×480
- Original goal: Interact with games in full 3D
- Computer vision happy with real-time depth and image
 - Games, HCI, etc
 - Action recognition
 - Image based modelling of dynamic scenes
- Fastest selling electronic appliance ever!!
- Other products that use PrimeSense sensor

