

Assignment 2 AI

Team - 25

- Pranay Gupta, 20161088
- Shubh Maheshwari, 20161170

$R(s, a)$ = reward for states all states

-25	0	0	0
0	0	0	0
2.5	0	0	0
0	-2.5	0	25

Step Reward = -5

$U_t(i,j)$ = Utility of state i,j at time t

Gamma = 0.9

Terminal state = [(0,0),(3,3)]

Wall = [(1,2)]

Value Iteration

We run this algorithm for all states except terminal states and walls (i, j) denotes i^{th} and j^{th} column until we reach tolerance

Iteration 1 :

Epoch:0 Error:15.000000

Updating:0 1 -5.000000

Updating:0 2 -5.000000

Updating:0 3 -5.000000

Updating:1 0 -5.000000

Updating:1 1 -5.000000

Updating:1 3 -5.000000

Updating:2 0 -2.500000

Updating:2 1 -5.000000

Updating:2 2 -5.000000

Updating:2 3 15.000000

Updating:3 0 -5.000000

Updating:3 1 -7.500000

Updating:3 2 15.000000

$V > V$

V V - V
V < V V
^ > > ^
-25.000 -5.000 -5.000 -5.000
-5.000 -5.000 0.000 -5.000
-2.500 -5.000 -5.000 15.000
-5.000 -7.500 15.000 25.000

This is the first iteration and so basically the policy can be seen as a randomized version where there are many issues with it and clearly we won't be able to reach the end state.

Iteration 2 :

Epoch:1 Error:13.000000

Updating:0 1 -10.000000
Updating:0 2 -10.000000
Updating:0 3 -10.000000
Updating:1 0 -8.000000
Updating:1 1 -10.000000
Updating:1 3 6.000000
Updating:2 0 -7.250000
Updating:2 1 -8.250000
Updating:2 2 8.000000
Updating:2 3 16.000000
Updating:3 0 -8.250000
Updating:3 1 3.250000
Updating:3 2 16.000000

V V > V
V < - V
^ > V V
> > > ^

-25.000 -10.000 -10.000 -10.000
-8.000 -10.000 0.000 6.000
-7.250 -8.250 8.000 16.000
-8.250 3.250 16.000 25.000

If we see in the first iteration the policy had a lot of problems. For example we would be stuck in this loop (see position (3, 0) and (2, 0)) of going up and down. So we would certainly need to

update our policy. This is the second iteration so it's basically just making the initial policy changes which are deemed good and which would result in us reaching the terminal state from most states. But still there are some issues with this. For ex we see the position (1,0) and (2,0). We would get stuck in this loop.

Iteration 3 :

Epoch:2 Error:8.975000

Updating:0 1 -15.000000
 Updating:0 2 -15.000000
 Updating:0 3 -2.200000
 Updating:1 0 -12.600000
 Updating:1 1 -13.225000
 Updating:1 3 9.000001
 Updating:2 0 -10.450000
 Updating:2 1 0.725000
 Updating:2 2 10.200001
 Updating:2 3 17.400000
 Updating:3 0 -3.950000
 Updating:3 1 4.800000
 Updating:3 2 17.400000

V V > V
 V **V** - V
 > > V V
 > > > ^

-25.000 -15.000 -15.000 -2.200
 -12.600 -13.225 0.000 9.000
 -10.450 0.725 10.200 17.400
 -3.950 4.800 17.400 25.000

This iteration finally makes the desired changes and we see that we would reach the terminal state if we start from any position.

Iteration 8 :

Epoch:7 Error:1.365857

Updating:0 1 -8.519598
 Updating:0 2 -2.177538

Updating:0 3 4.522495
Updating:1 0 -5.295281
Updating:1 1 -1.530580
Updating:1 3 11.776512
Updating:2 0 1.342914
Updating:2 1 5.555642
Updating:2 2 12.478209
Updating:2 3 18.048256
Updating:3 0 1.786273
Updating:3 1 8.289510
Updating:3 2 18.048256

V V > V
V V - V
> > V V
> > > ^

-25.000 -8.520 -2.178 4.522
-5.295 -1.531 0.000 11.777
1.343 5.556 12.478 18.048
1.786 8.290 18.048 25.000

Iteration 13 :

Epoch:12 Error:0.047148

Updating:0 1 -7.007732
Updating:0 2 -1.492260
Updating:0 3 4.768949
Updating:1 0 -3.999558
Updating:1 1 -0.908652
Updating:1 3 11.805258
Updating:2 0 1.890977
Updating:2 1 5.741687
Updating:2 2 12.499777
Updating:2 3 18.055481
Updating:3 0 2.074980
Updating:3 1 8.353044
Updating:3 2 18.055481

V V > V
V V - V

> > V V

> > > <

-25.000 -7.008 -1.492 4.769

-4.000 -0.909 0.000 11.805

1.891 5.742 12.500 18.055

2.075 8.353 18.055 25.000

Iteration 14 :

Updating:0 1 -8.839069

Updating:0 2 -4.432712

Updating:0 3 1.202816

Updating:1 0 -6.233100

Updating:1 1 -3.887096

Updating:1 3 8.324695

Updating:2 0 -1.138901

Updating:2 1 2.117785

Updating:2 2 9.057183

Updating:2 3 15.730927

Updating:3 0 -1.130982

Updating:3 1 4.963507

Updating:3 2 15.730927

Epoch:13 Error:0.005090

V V > V

V V - V

> > V V

> > > ^

-25.000 -8.839 -4.433 1.203

-6.233 -3.887 0.000 8.325

-1.139 2.118 9.057 15.731

-1.131 4.964 15.731 25.000

Iteration 29(Final State)

Updating:0 1 -6.965829

Updating:0 2 -1.475695

Updating:0 3 4.774305

Updating:1 0 -3.956245

Updating:1 1 -0.886900
 Updating:1 3 11.805555
 Updating:2 0 1.910087
 Updating:2 1 5.746768
 Updating:2 2 12.500000
 Updating:2 3 18.055555
 Updating:3 0 2.082969
 Updating:3 1 8.354579
 Updating:3 2 18.055555
 Epoch:29 Error:0.000000
 V V > V
 V V - V
 > > V V
 > > > ^

-25.000 -6.966 -1.476 4.774
 -3.956 -0.887 0.000 11.806
 1.910 5.747 12.500 18.056
 2.083 8.355 18.056 25.000

Result : We have finally reached the optimum policy state and as we can see we would reach the terminal state now if we start from any position on the grid.

Utility:

-25.000 -6.966 -1.476 4.774
 -3.956 -0.887 0.000 11.806
 1.910 5.747 12.500 18.056
 2.083 8.355 18.056 25.000

Final Policy:

V V > V
 V V - V
 > > V V
 > > > ^

If we start from (3,0)

We will always take right (3,0) => (3,1) => (3,2) => (3,3)

If we start from (1,0)

We will still reach (3,3) but we will take +2.5 reward and avoid -2.5 reward
 (1,0)=>(2,0)=>(2,1)=>(2,2)=>(3,2)=>(3,3)

Linear Programming

Values of X

XXXXXXXXX X	States	X
(State,Actions)		
(0, 0)	noop	0
(0, 1)	North	0
(0, 1)	South	0
(0, 1)	East	0
(0, 1)	West	0
(0, 2)	North	0
(0, 2)	South	0
(0, 2)	East	0
(0, 2)	West	0
(0, 3)	North	0
(0, 3)	South	0
(0, 3)	East	0
(0, 3)	West	0
(1, 0)	North	0
(1, 0)	South	0.01699879 6
(1, 0)	East	0

(1, 0)	West	0
(1, 1)	North	0
(1, 1)	South	0.02687777 85
(1, 1)	East	0
(1, 1)	West	0
(1, 3)	North	0
(1, 3)	South	0
(1, 3)	East	0
(1, 3)	West	0
(2, 0)	North	0
(2, 0)	South	0
(2, 0)	East	0.12611138 55
(2, 0)	West	0
(2, 1)	North	0
(2, 1)	South	0
(2, 1)	East	0.22490121 02
(2, 1)	West	0
(2, 2)	North	0
(2, 2)	South	0
(2, 2)	East	0.33880107 57
(2, 2)	West	0
(2, 3)	North	0
(2, 3)	South	0.30115651

		17
(2, 3)	East	0
(2, 3)	West	0
(3, 0)	North	0
(3, 0)	South	0
(3, 0)	East	1.12512348 73
(3, 0)	West	0
(3, 1)	North	0
(3, 1)	South	0
(3, 1)	East	1.02509878 98
(3, 1)	West	0
(3, 2)	North	0
(3, 2)	South	0
(3, 2)	East	0.94884348 83
(3, 2)	West	0
(3, 3)	noop	1

Expected Reward from linear programming :

2.082968874

Expected Reward from Value Iteration :

2.083

Why is the reward similar?

We try to maximize the utility/reward in both the methods of solving the MDP, so they would both end up achieving the same result if we try make them as accurate as possible. In VI, the

reward in the start state is the utility of selecting the best paths possible to the terminal states. In LP, the paths we get match the ones in VI so they will also consider similar probabilities. Now the reward in LP is the summation of the reward \times for each state,action pair. This will correspond to the value we get in VI if we assume a small delta, since a large delta would not allow enough iterations so that our VI spreads out enough and approximates the utilities of different states enough times. So if we use a delta not near 0, the values in VI and LP might not match but on using delta near 0, as in our case the rewards match in both of them.